

SRIVIN PLATFORMS

Delivery partner to RacingAhead / SportSeek

SportSeek — Unified Sports Ecosystem Platform

Project Delivery Model

How Srivin Platforms delivers SportSeek Phase 2A with an AI-native pod: public launch on 20 October 2026, then two features in production every month

Document control	
Document	SRV-SPS-2A-DLM-002
Programme	SportSeek Phase 2 — Phase 2A (Harden & Unify)
Client	RacingAhead / SportSeek, Hyderabad
Prepared by	Srivin Platforms — Technical Programme Management Office
Version / date	v2.0 (AI-native release train) — 27 September 2026
Status	Issued for client review — supersedes v1.0
Reference baseline	SportSeek Phase 2 SOW v1.0 (15 Jul 2026); SportsSeek App High Level Requirements
Classification	Confidential — RacingAhead / SportSeek and Srivin Platforms only

Contents

1. Purpose and scope
2. Our understanding of the engagement
3. Why AI-native, and what it changes
4. Delivery model at a glance
5. The AI-augmented delivery lifecycle
6. R1: the 20 October launch
7. The monthly release train
8. The AI-native pod
9. RACI matrix
10. Governance, escalation and decisions
11. Communication plan
12. Quality management
13. Environments, releases and feature flags
14. AI governance: security, confidentiality and IP
15. Knowledge transfer and transition
16. Change management
17. Risk management
18. Hypercare, warranty and support
19. Delivery metrics
20. Tooling
21. Commercial alignment
22. Our commitments to SportSeek

1. Purpose and scope

This document sets out how Srivin Platforms will deliver Phase 2A of the SportSeek programme for RacingAhead / SportSeek. It covers the AI-augmented lifecycle, the monthly release train, the pod, governance, the guardrails around AI, quality, knowledge transfer and support. Every ceremony, gate and control described here is one we will run.

The schedule, feature catalogue and resource loading are in SRV-SPS-2A-PLN-001 (WBS & Integrated Project Plan). The engineering approach is in SRV-SPS-2A-TAD-003 (Technical Approach). Scope references cite the SportSeek Phase 2 SOW v1.0 as "SOW §n".

2. Our understanding of the engagement

Phase 1 is live with a User App, a Partner App and an Admin Web Portal, built on **React Native, .NET and PostgreSQL**, with a payment gateway and SMS alerts already integrated (SOW §2.1). Phase 2A must unify identity, generalise partners, harden bookings, add a Transaction Ledger, expand notifications and deliver a web application (SOW §3.1). It must do so without disrupting a live product, and with SportSeek owning every account, repository and tool (SOW §5, §15).

SportSeek chose Srivin for **speed to market**. A plan that simply staffs a large team to a conventional timeline would give SportSeek nothing it could not do in-house. This model is built to put working features in customers' hands every month, starting on launch day.

3. Why AI-native, and what it changes

Conventional delivery	Srivin AI-native delivery
Large team; much of the effort goes into writing first-draft code, tests and documents	Small senior pod; AI agents produce first drafts, and engineers spend their time on design, review and hard problems
Quarterly-style milestones; value arrives late	A release every month, two features each , from 20 Oct 2026
Documentation written at the end, often stale	Documentation generated and updated with every change; always current
KT as a 2–4 week event after each milestone	Continuous KT: SportSeek developers work inside the pod; a signed checklist every release
Regression testing grows slower as scope grows	AI-generated regression suites grow with every feature and run on every merge

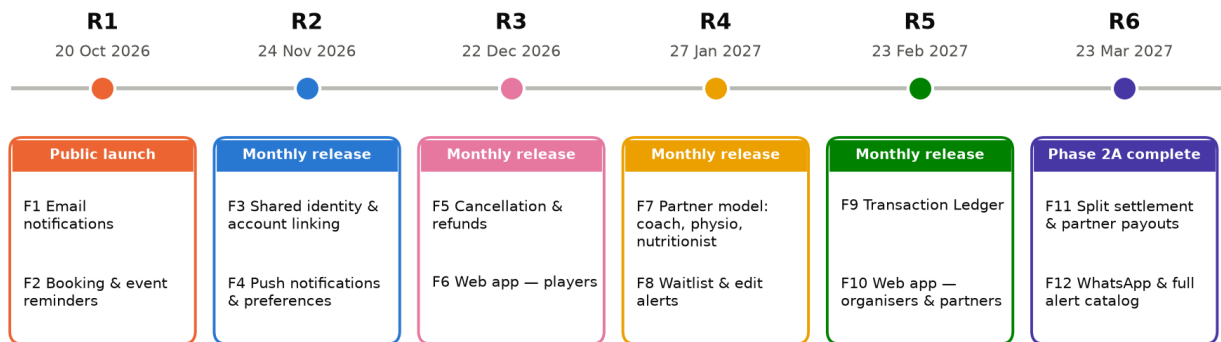
3.1 Principles

Principle	In practice
Humans own every outcome	AI drafts; a named senior engineer reviews, tests and merges. No AI output reaches production without human review and passing quality gates.
Specs before code	Every feature starts from an approved specification (story, API contract, data changes, UI, acceptance tests). The spec is what the agents build from, so quality starts upstream.
Fixed dates, flexible scope	The monthly train always leaves. A feature that is not ready rides the next one behind a feature flag.
Production first	SportSeek is live with real payments: regression, rollback and monitoring are in place before anything ships.
Transparent by default	SportSeek sees the same backlog, demos, metrics and AI usage data we do.
SportSeek ownership	All code, prompts/specs, documentation and AI-generated artefacts are SportSeek's (SOW §19). Everything lives in SportSeek's accounts.

What AI does not speed up, and how we plan for it

SportSeek decisions, UI approvals, UAT, app-store review, manual penetration testing and third-party approvals (Razorpay Route, DLT and WhatsApp templates) take the time they take. The plan gives each one a dated slot and a need-by date. When one is late, the affected feature moves to the next train and the release still ships.

4. Delivery model at a glance



Fixed dates, flexible scope: a feature that is not release-ready rides the next train behind a feature flag. It never delays the train.

Figure 1 — Six monthly releases, twelve features

Layer	Mechanism	Cadence	Who decides
Programme	Release-train plan, RAID, dependency register, change control	Monthly release review	Sponsor + Product Owner
Release	Gates: Specs approved → Feature complete → UAT sign-off → Go/No-Go → KT signed	Every month	Product Owner
Build	Weekly cycles inside each train; daily stand-up; weekly demo	Weekly	Product Owner (priority), pod (how)
Launch (R1)	War-room with daily triage to 20 Oct	Daily	Launch Lead + Product Owner

5. The AI-augmented delivery lifecycle

AI is used in every stage, and every stage has a human gate. This table is the core of the model.

Stage	What AI does	What people do	Human gate
Understand	Reads the Phase 1 codebase: architecture map, ERD, API catalogue, dependency and vulnerability scan, defect clustering	Architect validates findings and judges risk	Assessment accepted at Scope Gate
Specify	Drafts stories, Given/When/Then criteria, edge cases, API contracts and test cases from the SOW and workshops	Analyst and PO challenge, refine and prioritise	PO approves spec (Definition of Ready)
Design	Explores layouts and copy within the design system; exports tokens to code	Designer leads; usability checks with users	SportSeek UI approval (2 working days)
Build	Coding agents implement from the spec and generate unit tests, following repo conventions and module boundaries	Engineers review, refactor, handle complex logic and own the merge	Human code review + CI gates
Verify	Generates API, regression and E2E tests and test data; first-pass AI code review; SAST triage	QA curates suites and runs exploratory testing; specialists run manual pen tests	Feature complete; no open critical/high
Release	Drafts release notes, runbook updates and change summaries	DevOps releases; SportSeek runs UAT and Go/No-Go	UAT sign-off; Go/No-Go
Operate & hand over	Keeps docs current with every merge; answers questions over the codebase; triages logs and alerts	On-call engineers resolve; SportSeek team takes ownership	KT checklist signed

5.1 Human-in-the-loop rules (non-negotiable)

1. No AI-generated change merges without review and approval by a named engineer, who is accountable for it as if they wrote it.
2. All quality gates apply equally to AI-generated and hand-written code: tests, coverage, static analysis, security scans.
3. Security-sensitive areas (identity, payments, refunds, payouts, KYC data) get a second senior reviewer and a targeted manual pen test before release.
4. Agents work only in non-production environments with synthetic or masked data; they have no production credentials.
5. Generated dependencies and snippets are licence-checked; no code is accepted that the reviewing engineer cannot explain.

6. R1: the 20 October launch

R1 takes SportSeek public on a hardened, security-tested and monitored Phase 1 product, with **F1 Email notifications** and **F2 Booking & event reminders** on day one.

Time (IST)	Activity	Participants	Output
09:30 daily	Launch stand-up	Launch pod, SportSeek PO / tech lead	Launch board updated; blockers owned
17:00 daily	Defect and risk triage	Launch Lead, QA, PO	Fix / defer / de-scope decisions logged
01 Oct	Launch Scope Gate	SportSeek PO & sponsor; Srivin lead & architect	Frozen scope incl. F1/F2; go-live criteria
14 Oct	Code freeze	Launch pod	Release candidate; store submissions
19 Oct	Go/No-Go	SportSeek sponsor (decides), PO, Srivin lead	Signed Go/No-Go record
20 Oct	Launch and war-room	Pod on-call; SportSeek ops	Live; hourly health reports on launch day

6.1 Go-live criteria (proposed for the 1 Oct Scope Gate)

1. No open severity-1/2 defects in launch journeys; AI-generated regression suite green.
2. VAPT of the launch surface with no open critical/high findings (SOW §6.2).
3. Load test passed at SportSeek's launch concurrency (D5).
4. Monitoring, alerting and on-call live; rollback rehearsed; backup restore verified.
5. Privacy policy, terms and DPDP notice published; store listings approved or in final review.
6. UAT sign-off by SportSeek testers. F1/F2 are included only if they meet the same bar; otherwise they follow within two weeks.

7. The monthly release train

Week	Pod	SportSeek	Gate
Week 1	Plan from approved specs; agents implement; daily review; next train's specs start	PO clarifications	Specs approved
Week 2	Build and integrate on Staging behind flags; weekly demo	Demo feedback	—
Week 3	Feature complete Friday; regression, security scans, targeted pen test	UI approvals for next train	Feature complete
Week 4	Hardening; UAT support; store submission	UAT (3 days); Go/No-Go Monday	UAT; Go/No-Go
Release week	Release Tuesday; one week of hypercare	KT checklist sign-off	KT signed

7.1 Ceremonies

Ceremony	When	Time-box	Participants	Output
Train planning	Monday, week 1	1.5 h	Pod, PO	Train goal: the two features, with acceptance criteria
Daily stand-up	Daily 09:30	15 min	Pod incl. SportSeek developers	Blockers raised and owned
Weekly demo	Every Friday	45 min	Pod, PO, SportSeek stakeholders	Working software shown; feedback into backlog
Spec & design review	Weekly (Thu)	1 h	Analyst, designer, architect, PO	Next train's specs and UI approved
Release review / retro	After each release	1 h	Pod, PO, sponsor	Metrics, lessons, AI-playbook improvements

7.2 Definition of Ready and Done

	Criteria
Ready	Spec approved: story, acceptance criteria, API contract, data changes, UI approved, dependencies cleared, test data known. If it is not ready, the agents do not start.
Done (feature)	Merged after human review; unit/API/E2E tests green; coverage and security gates passed; docs and API collection updated automatically and checked; behind a feature flag; demoed and accepted by PO.
Done (release)	Full regression green; targeted pen test where applicable; UAT signed; release notes and runbook updated; deployed via the pipeline into SportSeek infrastructure (SOW §14).

8. The AI-native pod

One cross-functional pod of about 10–11 people, all senior enough to judge AI output. Total planned effort is 67.25 person-months, against 154 in our conventional estimate for the same scope. Roles, skills and monthly loading are in the WBS & Integrated Project Plan, Section 14.

Role	Accountability in the AI-native model
Delivery Lead / TPM	Runs the release train; owns plan, RAID, dependencies, decisions log and AI governance; single point of accountability.
Solution Architect & AI Engineering Lead	Architecture and ADRs; owns the AI engineering playbook (spec templates, context and prompt standards, agent guardrails, review rules); second reviewer on security-sensitive code.
Product Analyst	Turns SOW and workshops into approved specs with AI assistance; UAT scenarios; traceability.
Product Designer	Design system and feature UI; usability checks; design tokens shared by mobile and web.
Senior .NET / .NET & PostgreSQL Engineers	Direct coding agents from specs; review, harden and own back-end and database changes; migrations and reconciliation.
Senior React Native / React Engineers	Direct agents on mobile and web; own device testing, performance, accessibility and store releases.
QA Automation Engineer	Curates AI-generated test suites; owns quality gates; exploratory testing; UAT support.
DevOps / SRE	Pipelines, environments, observability, releases, on-call.
Security & Performance Engineers (on demand)	Targeted pen tests (R2, R3, R6) and full VAPT (R6); load tests (R1, R5, R6).

8.1 SportSeek inside the pod

- **2 SportSeek developers recommended** (D9) join from R2 as full pod members, working with the same AI tooling. They learn the AI-native way of working as well as the code.
- **Product Owner** approves specs and accepts features weekly; **UAT testers** test in the hardening week of each release.
- By R6, SportSeek engineers have built and shipped parts of every module and can run the train themselves or scale it into Phase 2B.

9. RACI matrix

R = Responsible, A = Accountable, C = Consulted, I = Informed. Aligned with SOW §15.

Activity	SportSeek PO	SportSeek Tech / Ops	SportSeek UAT	Srivin Delivery Lead	Srivin Architect	Srivin Pod
Product priorities & release scope	A/R	C	I	C	C	I
Feature specs	A (approve)	C	I	C	R	R
UI/UX design	A (approve)	C	I	I	C	R
HLD / ADRs	C	A (approve)	I	I	R	C
AI usage policy & tooling approval	A	C	I	R	R	I
Development incl. AI-generated code & review	I	C (review)	I	A	C	R
Testing (functional, regression, security)	I	I	I	A	C	R
UAT execution & sign-off	A	C	R	C	I	C
Go/No-Go	A	C	C	R (prepare)	C	I
Third-party contracts & approvals	A/R	C	I	C	C	I
Infrastructure operation (until handover)	I	C	I	A	C	R
KT checklist per release	A	R (receive)	I	R	R	R
Change requests	A (approve)	C	I	R (assess)	C	I

10. Governance, escalation and decisions

Forum	Frequency	Attendees	Purpose
Release review (acts as steering)	Monthly, after each release	Sponsor, PO; Srvin Engagement Director, Delivery Lead	Release outcome, next train scope, metrics, top risks, change requests
Weekly demo & status	Every Friday	PO, SportSeek tech, pod	Working software; train on-track status; decisions log; dependency register
Design authority	Fortnightly / on demand	Architect, SportSeek tech reviewer	ADRs and technical decisions
Launch war-room	Daily to 6 Nov 2026	Launch pod, PO	Section 6

Level	Escalated to	Trigger	Response
L1	Engineer ↔ PO	Clarification or blocker	Same working day
L2	Delivery Lead ↔ PO / tech lead	Decision or dependency at risk of missing need-by date	1 working day
L3	Engagement Director ↔ Sponsor	A feature at risk of missing two trains; commercial issue	2 working days
Launch	Launch Lead → Delivery Lead → Sponsor	Any risk to 20 Oct or go-live criteria	4 business hours

Every SportSeek decision is logged with options, Srvin's recommendation and a need-by date derived from the train. A decision missing its date moves the affected feature, not the release.

11. Communication plan

Communication	Audience	Frequency	Owner
Weekly status (train progress, decisions, dependencies, RAID) — AI-drafted from tool data, reviewed by the Delivery Lead	PO, sponsor	Weekly (Fri)	Delivery Lead
Weekly demo of working software	SportSeek stakeholders	Weekly	Pod
Release notes (business and technical)	SportSeek business, support, ops	Monthly	QA / Delivery Lead
Release review pack with delivery and AI metrics	Sponsor, PO	Monthly	Delivery Lead
Incident communication	SportSeek ops, PO, sponsor (S1)	As needed	On-call lead

12. Quality management

Level	Scope	How AI helps	When
Unit	Domain logic and edge cases	Generated with code from the spec; reviewed	Every commit
API / integration	Endpoints, DB (Testcontainers), adapters	Generated from OpenAPI contracts	Every merge
Regression	All live journeys, mobile and web	Suite generated from journey inventory; grows each release	Nightly + pre-release
Security	SAST/SCA/secrets in CI; DAST; manual pen tests	AI triage of scanner findings	CI + R1, R2, R3, R6
Performance	Load at SportSeek targets (D20)	Scripts generated from API catalogue	R1, R5, R6
UAT	Business scenarios	Scenario scripts drafted from specs	Hardening week

Severity	Definition	Example
S1 Critical	Production down, data loss, security breach, or payments failing broadly	Payments not captured
S2 High	Core journey broken for some users, or a wrong financial or identity outcome	Duplicate identity on linking
S3 Medium	Feature misbehaves with a workaround	A ledger filter ignores one status
S4 Low	Cosmetic	Alignment or copy issues

13. Environments, releases and feature flags

- **Dev / Staging / Production** in SportSeek accounts. Agents and engineers work in Dev; Staging holds synthetic or masked data only; Production changes go only through the pipeline.
- **Continuous delivery behind flags**: back-end changes deploy when ready and stay dark until the release; mobile ships on the monthly train via store submission in week 4.
- **Rollback**: expand/contract database migrations and flag kill-switches make every release reversible.
- **Minimum-version enforcement** on the apps, so API changes (e.g. identity in R2) never strand old clients.

14. AI governance: security, confidentiality and IP

Control	How it works
Approved tooling only	AI tools are used only under enterprise/business terms that exclude training on customer code and data, and only after SportSeek approves them (D21). The tool list is part of the governance charter.
Data handling	No production personal, KYC or transaction data in prompts, agent contexts or non-production environments. Secrets are never placed in prompts; secret scanning runs in CI.
Least privilege for agents	Agents run in development workspaces with repository access only. They have no production credentials and cannot deploy; deployment is a pipeline action approved by a person.
Human accountability	Every merge has a named human reviewer. Security-sensitive changes have two.
IP and licensing	All outputs, including code, tests, specs, prompts and documentation, belong to SportSeek on payment (SOW §19). Dependencies are licence-scanned; the reviewing engineer confirms generated code contains no copied third-party material.
Audit trail	Pull requests record AI assistance, reviewers and checks. AI usage metrics are reported monthly.
DPDP	AI governance sits inside the same DPDP-aligned data controls as the rest of the engagement.

15. Knowledge transfer and transition

- **Continuous KT:** SportSeek developers build inside the pod from R2, and specs, ADRs, API docs and runbooks are generated and kept current with every merge. There is no end-of-milestone documentation cliff.
- **Per-release KT checklist** (SOW §12.2) signed within 5 working days of each go-live, covering walkthrough, docs, runbook, and a change made by a SportSeek engineer.
- **Codebase assistant for SportSeek:** we set up an AI assistant configured on SportSeek's own repositories and documentation, within SportSeek's approved tooling. SportSeek's team can keep asking "how does X work?" long after handover.
- **Infrastructure:** operate → co-operate (from R3) → handover, with the operations runbook delivered in R6 (SOW §15.1).
- We are proposing continuous KT instead of a 2-week-to-1-month KT window after each milestone (SOW §12.2). The checklist sign-off the SOW requires is retained.

16. Change management

SOW §18 applies: a written request, then a Srivin impact assessment within 5 business days, then SportSeek's decision, then an updated baseline. The release train makes change cheaper. New scope is slotted into a future train, and swapping features of similar size between trains is a PO decision recorded in the decision log, with no change request needed.

17. Risk management

#	Risk	P / I	Mitigation	Owner
R1	AI tooling not approved by SportSeek on day one	L-M / H	Enterprise-terms tools; written AI policy at kick-off; fallback plan at conventional pace	Delivery Lead
R2	Launch blockers found late in Phase 1	M / H	AI-assisted assessment in 3 days; Scope Gate; F1/F2 de-scoped before launch is risked	Delivery Lead
R3	SportSeek decision/approval pace lags the monthly train	M / H	1-day decision and 2-day UI approval agreement; weekly decision log; affected feature moves, not the release	SportSeek PO
R4	Quality or security issues in AI-generated code	L-M / H	Human review, dual review on sensitive code, full gates, targeted pen tests	Architect
R5	Over-reliance on AI erodes team understanding	L / M	Engineers must explain every merge; SportSeek developers in the pod; living docs	Architect
R6	Third-party approvals (Route, DLT, WhatsApp)	M / M	Need-by dates months ahead; those features in the last train	SportSeek
R7	Duplicate identities in existing data	M / M	Early duplicate report; approved merge rules; rehearsed migration	Architect
R8	App-store review delays a mobile release	L-M / M	Submission in week 4; web and back end unaffected; features stay behind flags	Delivery Lead

18. Hypercare, warranty and support

- **Hypercare:** 20 Oct – 6 Nov 2026 for the launch; one week after every later release; 23 Mar – 2 Apr 2027 for R6.
- **Warranty:** 90 days per release at no additional cost (SOW §16). **Phase 1 support** included (SOW §16).
- **AI-assisted operations:** log and alert triage speeds up diagnosis; fixes follow the same review and release gates.

Severity	Acknowledge	Workaround / restore	Permanent fix
S1 Critical	30 min (24×7 in hypercare)	4 hours	Next hotfix
S2 High	2 business hours	1 business day	Next release or hotfix
S3 Medium	1 business day	As agreed	Next release
S4 Low	2 business days	—	Backlog

Note

These SLA figures are a Srivin proposal. The SOW leaves them open (SOW §16 ▀, D16).

19. Delivery metrics

Area	Metric	Target
Throughput	Features in production per month	2 per month from R1
Predictability	Releases on their fixed date; features that moved trains	100% on date; moved features reported with reasons
Speed	Lead time from approved spec to production; deployment frequency	Reported monthly
Quality	Escaped S1/S2 defects; change failure rate; mean time to restore	Zero S1/S2 escapes as the aim
Security	Open critical/high vulnerabilities at release	Zero
AI effectiveness	Share of merged changes that were AI-assisted; review time; rework rate on AI-generated code	Reported monthly; used to tune the playbook
Ownership	KT checklist items signed; SportSeek-authored merges	Rising every release

20. Tooling

Purpose	Tool	Ownership
AI engineering	AI coding agents (e.g. Claude Code) and assistants under enterprise terms, approved by SportSeek	Configured on SportSeek repositories; usage logged
Backlog, specs, tests	JIRA or Trello (SportSeek decision, D8)	SportSeek-owned; Srivin operates
Code & CI/CD	SportSeek version control and pipelines	SportSeek-owned
Design	Figma or tool agreed with SportSeek	Files handed over (SOW §13)
API docs	OpenAPI + Postman collections, generated and checked each merge (SOW §13)	In SportSeek repositories
Observability	Logging, error tracking, metrics, uptime (SOW §6.2)	SportSeek cloud

21. Commercial alignment

The release train also offers a simpler commercial shape, for inclusion in the proposal the SOW requests:

- **Price per release:** a fixed fee for each monthly release (two features), invoiced on UAT sign-off. SportSeek pays for delivered, accepted features and can approve costs one release at a time, as the SOW Next Steps ask.
- **Launch (R1)** priced separately as a fixed-scope hardening engagement.
- **Optional continuation** into Phase 2B on the same train from April 2027, at the same per-release structure.
- Rates, payment terms and any outcome-linked elements will be in the commercial proposal.

22. Our commitments to SportSeek

1. **Public on 20 Oct 2026**, with two new features on day one.
2. **Two features in production every month**, on fixed release dates, to Phase 2A completion on **23 Mar 2027**.
3. **A lean senior pod** (67.25 person-months, less than half the conventional estimate), with every AI output reviewed and owned by a named engineer.
4. **Transparency**: weekly working software, monthly metrics including AI usage.
5. **Ownership**: SportSeek's accounts, code, docs and AI artefacts; SportSeek developers able to run the train themselves.